

Ideology as topic correlation structure: Inference method and example application to carbon tax public opinion

Journal:	<i>Environmental Data Science</i>
Manuscript ID	EDS-2023-0007.R1
Manuscript Type:	Application Paper
Date Submitted by the Author:	n/a
Complete List of Authors:	Puelma Touzel, Maximilian; Universite de Montreal, Computer Science and Operations Research Lachapelle, Erick; Universite de Montreal, Political Science
Keywords:	Machine Learning, Data Science, Policy < Climate Change, Artificial Intelligence, Public Opinion
Abstract:	<p>Ideology is a system of ideas that helps people rationalize their beliefs and explain the world. As an object of study, ideology is amenable to a generative modelling approach within the text-as-data paradigm. Here, we characterize ideology as the relational structure among topics. To access this latent structure, we use the Structural Topic Model to infer topics and the weights with which individual opinions mix topics. We fit the model to survey data of open-ended responses of Canadians elaborating on their support of or opposition to a carbon tax. We demonstrate that the latent topic mixture, in contrast to the observed word usage, is an efficient representation for the machine classification of responses, allowing for near-perfect accuracy in discriminating opposition from support. We operationalize relational structure as the correlation structure of the topic mixture statistics over a surveyed group. We evaluate four principled low-order statistics with intuitive interpretations for ideology: specificity, variability, expressivity, and alignment. We find higher specificity and alignment, much less expressivity, and similar variability in the ideology behind oppose responses relative to support responses. We discuss the implications of the results for climate policy debates and of our approach to analyzing ideology for computational social science.</p>

SCHOLARONE™
Manuscripts



1

201 RESEARCH ARTICLE

302

403

504

605

Ideology as topic correlation structure: Inference method and example application to carbon tax public opinion

109 Maximilian Puelma Touzel^{1,2} * and Erick Lachapelle³
110

121 ¹Mila, Québec AI Institute, Montréal, Québec, Canada

122 ²Department of Computer Science and Operations Research, Université de Montréal, Montréal, Québec, Canada

123 ³Department of Political Science, Université de Montréal, Montréal, Québec, Canada

124 *Corresponding author. Email: max.puelmatouzel@gmail.com

125 Received: 31 January 2023; Revised: In revision; Accepted: Pending revision

126 Keywords: Machine Learning, Data Science, Climate Change Policy, Artificial Intelligence, Public Opinion

127 Abstract

128 Ideology is a system of ideas that helps people rationalize their beliefs and explain the world. As an object of
129 study, ideology is amenable to a generative modelling approach within the text-as-data paradigm. Here, we charac-
130 terize ideology as the relational structure among topics. To access this latent structure, we use the Structural Topic
131 Model to infer topics and the weights with which individual opinions mix topics. We fit the model to survey data
132 of open-ended responses of Canadians elaborating on their support of or opposition to a carbon tax. We demon-
133 strate that the latent topic mixture, in contrast to the observed word usage, is an efficient representation for the
134 machine classification of responses, allowing for near-perfect accuracy in discriminating opposition from support.
135 We operationalize relational structure as the correlation structure of the topic mixture statistics over a surveyed
136 group. We evaluate four principled low-order statistics with intuitive interpretations for ideology: specificity,
137 variability, expressivity, and alignment. We find higher specificity and alignment, much less expressivity, and similar
138 variability in the ideology behind oppose responses relative to support responses. We discuss the implications of
139 the results for climate policy debates and of our approach to analyzing ideology for computational social science.

321 Impact Statement

322 What people think about climate change and climate policy strongly depends on their politics. To better
323 understand the role played by ideology in shaping people's attitudes, we assess people's ideological thinking
324 by analyzing the ways in which they explain their stance in open-ended survey responses. We develop new
325 ways of measuring how topics are talked about together. Independent of what respondents discussed, we find
326 people elaborating on opposition to the carbon tax did so in a more focussed and structured way than did
327 those who support a carbon tax. This knowledge could be used to help design more effective climate policy
328 and communications.

443 1. Introduction

444 Personal values and prior beliefs can have an out-sized effect in shaping how individuals interpret the
445 world and how they respond when prompted to opine about specific aspects of it. One helpful construct
446 for understanding this cognitive and collective process is ideology, a system of interconnected beliefs
447 that is much discussed in the context of debates surrounding the role of political partisanship and the

2 *Puelma Touzel et al.*

1

201

302

403

504

605

706

807

908

109

110

121

132

143

154

165

176

187

198

209

220

232

243

254

265

276

287

298

309

321

332

343

354

365

376

387

398

409

421

432

443

454

465

476

487

498

509

510

521

53

54

55

causes of political polarization [Baldassarri and Gelman \(2008\)](#). However, how to analyze ideology as a correlation structure on beliefs is a methodological challenge yet unresolved [Kalmoe \(2020\)](#). In this study, we propose an approach to inferring ideological reasoning from open-ended survey data, and ask to what extent topic correlations can reveal differences in the way ideological groups rationalize their positions on carbon tax policy. Given the well-documented anti-tax rhetoric of conservative political discourse ([Lakoff, 2010](#)), we expect ideologically-driven differences in the topic structures called to mind among supporters and opponents of carbon tax policy.

There are many key societal challenges to which methods that reveal ideology could contribute. A prime example is mitigating the worst effects of climate change. Climate change mitigation is in many countries a hotly contested political issue that is characterized by extensive ideological and partisan conflict ([Birch, 2020](#); [McCright and Dunlap, 2011](#); [McCright et al., 2016](#)). Addressing this challenge requires significant government intervention to shift energy production and consumption away from CO₂-emitting fossil fuels ([McGlade and Ekins, 2015](#); [Welsby et al., 2021](#)). In this context, a price on carbon emissions is often proposed by economists as being simple, flexible, and easy to put on a rising schedule, encouraging emissions reductions across economic sectors at a low aggregate cost ([Baranzini et al., 2000](#); [Aldy and Stavins, 2012](#)). Policy design has addressed public resistance to carbon tax policy measures via the targeted use of funds in the form of lump sum rebates to taxpayers ([Klenert et al., 2018](#)). Despite these mechanisms, however, a lack of public acceptability remains a major obstacle to the adoption of substantive carbon taxes across countries ([Haites, 2018](#); [Lachapelle, 2017](#); [Rabe, 2018](#)).

Public support for carbon pricing is typically structured by ideology, with conservatives on the right of the political spectrum generally in opposition, and those on the left generally supporting the policy ([Drews and Van den Bergh, 2016](#)). This structuring of carbon tax policy attitudes by ideology is evident even in countries that have implemented tax-and-rebate schemes designed to build public support by providing rational, monetary incentives that offset policy costs. For instance, at current pricing levels in Canada (one of few countries that have implemented a tax-and-rebate scheme), 8 of 10 households are estimated to receive more than they pay as a result of the policy, and yet roughly 8 out 10 surveyed conservatives oppose the policy ([Mildenberger et al., 2022](#)). [Mildenberger et al. \(2022\)](#) probe that result using interventions and find conservatives underestimate the size of the rebate they receive. Such behaviour is not necessarily irrational: strong priors can make rational decision-making insensitive to new data ([Gershman, 2019](#)). Even after being provided with information on the size of their rebate, individuals may continue to believe that they are net losers. One explanation for these results is that a broader system of values—an ideology—is driving subjects' decision-making. This putative cognitive belief system would then also likely manifest when subjects justify their opposition to the tax. A better understanding of how ideology underlies carbon tax opinion could thus help inform the design of more effective carbon pricing policy in general, perhaps via more effective communication of the policy's benefits. This might involve first identifying, then appealing to (and in some instances directly challenging) issue frames commonly associated with carbon tax policy beliefs.

Issue frames, when written out, can be partially quantified by the semantic content of the writing. Quantitative semantic representations typically focus on frequency of word-use¹. However, single-word frequency measures do not expose how the same words can be used when talking about different things. Other widely used approaches, such as sentiment analysis, classify responses into only a few affective classes ('like'/'dislike'). By formulating a rich latent topic structure, *topic models* address both these limitations. Topic models are now an established approach to understanding human-generated responses in the social sciences ([Grimmer et al., 2022](#)). The Structural Topic Model in particular has been applied to understand open-ended responses on a carbon tax in Spain ([Savin et al., 2020](#)), Norway ([Tvinnereim et al., 2017](#)), and the US ([Povitkina et al., 2021](#)). Here, we make a similar application to data collected in Canada. Unlike these previous works, however, we focus on topic-topic correlations as a means to interrogate ideology. Our contribution is a set of measures and analyses that characterize this latent correlation structure as a means of inferring ideology from open-ended survey responses.

¹This includes more refined frequency measures such as TfIdf (see fig. 3).

1

201 To further motivate our application to carbon tax opinion, we hypothesized that opposition to carbon
302 taxes arises from a well-worn ‘tax is bad’ ideology, involving a handful of correlated ideas (‘distrust
403 in government’, ‘unfairness’, ‘infringement on personal freedom’, etc.) that mutually enforce each
504 other (Lakoff, 2010).

605 Here, we use the Canadian dataset collected by (Mildenberger et al., 2022), unique in the richness
706 of its meta-data, to fit the parameters of a generative bag-of-words model of word responses having a
807 latent structure of topics. We fit models to different response types: the subset of responses of those who
908 stated separately that they support the tax, of those who stated they opposed the tax, and of those two
109 groups combined. After validating the fitted models, we use it to infer topic structure conditioned on
110 the support and oppose response groups. We focus on the ways in which respondents mix topics when
121 supporting or opposing the carbon tax. We not only find that responses are highly discriminable using
132 these topic mixtures, but that there are clear differences in the correlation structure of the two response
143 types that supports our hypothesis.

144

145

176 2. Methods

187 2.1. Dataset

198

209 We analyzed a dataset of 3313 response documents of survey respondents from across Canada in
210 2019 on the question of support for the recently implemented carbon tax as published in Milden-
221 berger et al. (2022). In addition to asking respondents for a categorical opinion response in *sup-*
232 *port/opposition/notsure* (response type), demographic survey data were collected such as age, gender,
243 and province. Lifestyle data (car use, residence environment) and belief data (partisanship (measured
254 using party voting) and political ideology) were also collected. Ideology was measured with a uni-
265 dimensional scale asking respondents, “In politics, people sometimes talk of left and right. Where
276 would you place yourself on the following scale?”, with a 0 to 10 scale with extreme values labeled
287 “Far to the left” (0) and “Far to the right” (10). The mid-point was also labeled as “Centre” (5). In line
298 with research documenting a general trend in partisan dealignment (Dalton and Wattenberg, 2002),
309 partisanship was measured with a vote choice question asking respondents, “If a federal election were
330 held today, for which party would you vote for?” Reflecting the multiparty nature of Canadian politics,
321 response options included items for each of Canada’s main federal political parties: Liberal, Conser-
332 vative, New Democratic Party, People’s Party, and Bloc Québécois (for respondents in Quebec). An
343 option was also included for “I would not vote.” In the analysis, we recoded Ideology and Partisanship
354 into 3 categories each. See Mildenberger et al. (2022) for more details on this rich meta-data and see
365 the Supplemental Material for the specific reductions of categories that we used. In addition to cate-
376 gorical responses, a central open-ended question in the survey asked respondents to elaborate on why
387 they choose the categorical response *support/oppose/not sure*. For this analysis, we preprocessed these
398 open-ended responses. French responses were first translated into English using the Google Translate
409 API. The response corpus was pre-processed through automated spell-checking, removal of stop words,
440 and reduction to word stems. Words were tokenized: each word in the vocabulary was assigned an index
421 and each word sequence in each response by transformed to its vector representation whose sequence of
432 elements are the sequence of the corresponding word indices (e.g. the response {‘I’, ‘love’, ‘icecream’} might map to {3403, 5937, 2345}).

443 The *not sure* response types were analyzed in addition to the *support* and *oppose* types. They
454 were only 11% of respondents, the majority of which gave responses to the effect of “I don’t know
465 enough to answer” (see Supplemental Material). Thus, here we focus the presentation of our analysis
476 on contrasting the *support* and *oppose* types.

487

498

509

510

521

53

54

55

4 Puelma Touzel et al.

1 201 2.2. Word relevance

302 We used a standard term-relevance measure that combines term frequency (as a measure of term
403 relevance) and inverse document frequency (as a measure of term specificity):
504

- 605 • *term frequency*, $\text{tf}(v, d)$. The frequency of a vocabulary word v in a document d is
706 $\text{tf}(v, d) := \frac{n(v, d)}{n(d)}$, where $n(v, d)$ is the number of times term v is in document d and
807 $n(d) = \sum_{v \in \mathcal{V}} n(v, d)$ is the total number of words in document d . The term frequency over all
908 the documents is then,

$$109 \quad \text{tf}(v) := \frac{\sum_{d=1}^D n(d)\text{tf}(v, d)}{N_{\text{total}}},$$

110 where the denominator $N_{\text{total}} = \sum_{d=1}^D n(d)$ is just the total word count across all D considered
121 documents.

- 132 • *inverse document frequency*, $\text{idf}(v)$. Here,

$$143 \quad \text{idf}(v) = \frac{\log(D + 1)}{\log(n(v) + 1) + 1},$$

154 where $n(v)$ is the number of documents in which term v appears, i.e. $n(v, d) > 0$. Idf is like an
165 inverse document frequency.

198 Term frequency-inverse document frequency (Tfidf) for a word-document pair is then defined simply as
209 the product, $\text{Tfidf}(v, d) := \text{tf}(v, d)\text{idf}(v)$. This is then averaged over documents to create a word-only
210 measure. We used the *sklearn* package implementation of Tfidf that averages the normalized values to
221 remove dependence on the length of a document

$$226 \quad \text{Tfidf}(v) = \frac{1}{D} \sum_{d=1}^D \frac{\text{Tfidf}(v, d)}{\|\text{Tfidf}(\cdot, d)\|},$$

330 where $\|\mathbf{x}\| = \sqrt{\sum_{i=1}^N x_i^2}$ is the Euclidean norm (throughout the paper, we denote vectors with
321 bold symbols). As a word relevance metric, Tfidf aims to add more discriminability than frequency
332 alone by downweighting words that appear in many documents since these common words are less
343 discriminative. We computed word clouds and rank histograms using the Tfidf values.
354

365 2.3. Response-type Classification

376 We aimed to predict the response type label (*oppose/support*) for each response, first using the word-
387 level representation (*i.e.* which words appear and how many times), then using the inferred topic
398 mixture weights obtained from specific (e.g. K -dependent) fitted STMs on the combined dataset of
409 support and oppose responses (see the following sections for details on the topic model). As prepro-
410 cessing, we applied maximum absolute value rescaling to handle sparsity in the data for the word-level
421 representation, and standard rescaling in the topic mixture case. In both cases, we then performed logis-
432 tic regression. We then ranked features (words or topic weight vector components, respectively) by
443 their weighted effect size and then trained a classifier for each subset of ranked features of increasing
454 number, n . For each n , we ran 100 trials of a 2:1 train/test split and report the mean and variance of
465 classification accuracy.
476

487 2.4. The Structural Topic Model

498 Topic models are generative models that generate a set of words in a response document according to
509 some distribution. Topic models are typically bags of words models, which eschew grammar and syntax
510 typically found in language models.

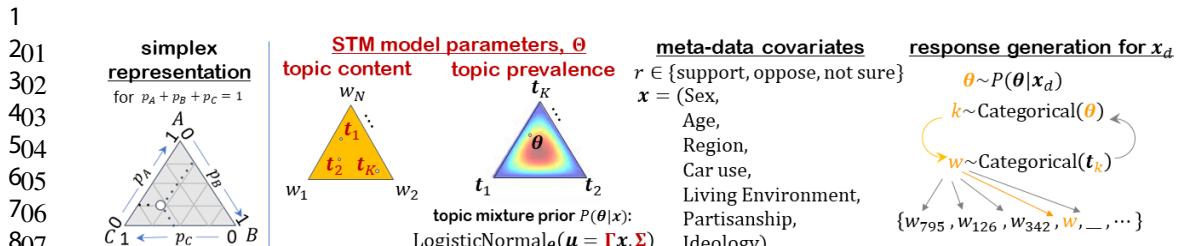


Figure 1. The Structural Topic Model for analysis of public opinion survey data. From left to right: Recall the simplex as the space of frequency variables. STM parameters define topic content through the topics' word frequencies and topic prevalence through parameters associated with the topic mixture prior. We use seven meta-data covariates. We also store the respondent's categorical response (in support, in opposition or not sure) to the issue question (here carbon tax). By sequentially sampling topics and words, the model produces a bag-of-words open-ended response elaborating on this opinion.

to focus only on the content and prevalence of words (*i.e.* sampling distributions do not condition on previously sampled words in a response). We select a topic model with rich latent structure: the *Structural Topic Model* (STM) (Roberts et al., 2014). Like the Correlated Topic Model (Blei and Lafferty, 2005), the STM uses a logistic normal distribution from which to sample the topic mixture weights on a document and can thereby exhibit arbitrary topic-topic covariance via the covariance parameter matrix of the logistic normal distribution. Unlike the Correlated Topic Model in which the mean and covariance matrix of the logistic normal distribution are learned parameters, the STM allows for parametrizing these using a linear model of the respondents' meta-data. We discuss our choice of meta-data model below. In the standard implementation of STM that we use here (fig. 1), only the mean of the logistic normal is made dependent on the meta-data. This adds flexibility beyond that offered by the CTM and makes the STM appropriate for datasets with non-trivial latent topic structure as we expect here. Specifically, our use of the STM model specifies the parameter tuple $\Theta = (\{t_k\}_{k=1}^K, \Sigma, \Gamma)$ as follows: (fig. 1):

1. **Topic content**: an underlying set of K topics indexed by k , $\{t_k\}_{k=1}^K$, where each topic, $t_k = (\beta_{k,1}, \dots, \beta_{k,|\mathcal{V}|})$, is a list of word frequencies on a given vocabulary \mathcal{V} indexed by v ($\sum_{v=1}^{|\mathcal{V}|} \beta_{k,v} = 1$ for all k), and
2. **Topic prevalence**: a prior distribution $P(\theta|x)$ on the topic mixture vector, $\theta = (\theta_1, \dots, \theta_K)$ (weights are normalized, $\sum_{k=1}^K \theta_k = 1$) that is conditioned on the meta-data class, x . Here, the distribution is chosen as a logistic normal with covariance matrix parameter Σ assumed to be the same across meta-data classes, and mean $\mu = \Gamma x$ is a linear transformation of x with matrix parameter Γ .

A single response sample for a given response document of length N words that comes from a respondent in meta-data class x is generated by sampling a topic mixture weight vector θ once from the weight vector distribution $P(\theta|x)$, then iteratively sampling a topic index, $z_n \in \{1, \dots, K\}$, from $\text{Categorical}_K(\theta)$ and a word, $w_n \in \mathcal{V}$ from the distribution formed from that topic's frequencies, $\text{Categorical}_{|\mathcal{V}|}(t_{k=z_n})$, N times to make the response $\{w_1, \dots, w_N\}$, with $n = 1, \dots, N$. Topic frequencies are represented in log-space according to the Sparse Additive Generative text model (Eisenstein et al., 2011). This topic content model allows for meta-data dependence. However, again confronted with the challenging model selection problem, and without strong hypotheses about how demographic-biased word use affect topic correlations (our primary interest), we chose to leave out this dependence.

1

201

302

403

504

605

706

807

908

109

110

121

132

143

154

165

176

187

198

209

220

232

243

254

265

276

287

298

309

330

321

332

343

354

365

376

387

398

409

440

421

432

443

454

465

476

487

498

509

510

521

53

54

55

2.5. STM Model Fitting

We built a Python interface to a well-established and computationally efficient STM inference suite written in the R programming language (Roberts et al., 2019). This STM package has a highly optimized parameter learning approach that largely overcomes the inefficiency of not using a conjugate prior². In particular, STM uses variational expectation maximization, with the expectation made tractable through a Laplace approximation. Accuracy and performance of the method is improved by integrating out the word-level topic and through a spectral initialization procedure. For details see (Roberts et al., 2014). The Python code that we developed to access this package is publicly accessible with the rest of the code we used for this paper in the associated GitHub repository.

The parameter fitting performed by this STM package uses a prior on its covariance matrix parameter Σ of the topic mixture weight prior that specifies a prior on topic correlations: $\sigma \in [0, 1]$ giving a uniform prior for $\sigma = 0$ and an independence (i.e. diagonal) prior for $\sigma = 1$ (see the package documentation for more details on σ). Unless otherwise stated, we used the default value of $\sigma = 0.6$. Note that our main analysis of statistical measures does not use Σ . The package also offers priors on Γ . We used the recommended default "Pooled" option. This uses a zero-center normal prior for each element, with a Half-Cauchy(1,1) prior on its variance.

We fit three model types based on what dataset they are trained on. Here, a dataset $\mathcal{D} = \{(r_d, w_d, x_d)\}_{d=1}^D$ is a set of (categorical response, open-ended response, respondent metadata) tuples. We fit one model to each of the support and oppose respondent data subsets, $\mathcal{D}_{\text{support}}$ and $\mathcal{D}_{\text{oppose}}$, as well as one to a combined dataset $\mathcal{D}_{\text{combined}} = \mathcal{D}_{\text{support}} \cup \mathcal{D}_{\text{oppose}}$ that joins these two subsets of tuples into one. The model parameters obtained from maximum likelihood fitting to a dataset are denoted $\hat{\Theta}$. As with other topic models, an STM takes the number of topics, K , as a fixed parameter (on which most quantities depend so we omit it to simplify notation). We thus obtained three K -dependent families of model parameter settings $\hat{\Theta}_m$, for $m \in \{\text{support, oppose, combined}\}$. With model parameters in hand, we obtained our object of primary interest: the models' topic mixture posterior distributions conditioned on response type $r' \in \{\text{support, oppose}\}$,

$$P(\theta|\Theta = \hat{\Theta}_m, r') = \sum_x P_m(\theta|x)p(x|r') \quad (1)$$

$$\approx \frac{1}{|\mathcal{D}_{r'}|} \sum_{d \in \mathcal{D}_{r'}} \text{LogisticNormal}(\Gamma_m x_d, \Sigma_m), \quad (2)$$

where we have substituted the Logistic Normal distribution for $P_m(\theta|x)$ and use the empirical average over x . To compare the ideologies behind support and oppose responses, we make the two comparisons from evaluating eq. (1):

$$P(\theta|\Theta = \hat{\Theta}_{\text{support}}, \text{support}) \xleftrightarrow{?} P(\theta|\Theta = \hat{\Theta}_{\text{oppose}}, \text{oppose}) \text{ and} \quad (3)$$

$$P(\theta|\Theta = \hat{\Theta}_{\text{combined}}, \text{support}) \xleftrightarrow{?} P(\theta|\Theta = \hat{\Theta}_{\text{combined}}, \text{oppose}). \quad (4)$$

Since we are interested in topic correlations, we do not need to derive effect sizes for particular meta-data covariates, which in general depend on the particular meta-data model, here given by which elements of Γ in $\mu = \Gamma x$ are non-zero and their sign. Consequently, we can sidestep the challenging meta-data model selection problem (Wysocki et al., 2022) by simply including all the measured covariates that might contribute (*i.e.* all elements of Γ are specified as real-valued free parameters). We thus selected a broad subset of the meta-data from Mildenberger et al. (2022): $x = (\text{age}, \text{sex}, \text{region}, \text{car use}, \text{partisanship}, \text{ideology})$.

²The widely used Latent Dirichlet Allocation topic model uses a Dirichlet prior and the variational Bayes algorithm for this reason, but lacks expressiveness as a result. Cambridge University Press

1

201 We analyzed results across a range of K to ensure validity of our conclusions. Models with larger
 302 K typically give higher model likelihood but are at higher risk of over-fitting. The model likelihood
 403 on a heldout test dataset will however exhibit a maximum likelihood at finite K as overfitting to the
 504 remaining ‘train’ samples becomes significant with increasing K . The value of K where the heldout
 605 likelihood achieves its maximum will depend on the relative size of the test data relative to the train
 706 data in the train-test split. Nevertheless, a loose upper-bound on K is the value at which the maximum
 807 of the split realization-averaged held-out likelihood for a train-biased train-test data split (here 9:1) is
 908 obtained. For the support, oppose, and combined response type models, maximums were observed at
 109 $K = 15, 18,$ and 20 , respectively for averages over 100 and 200 train-test split realizations for separate
 110 and combined response type models, respectively (see fig. 6e,j). Note that the strong subadditivity
 121 ($20/(15+18)=0.61$) in these data-selected topic numbers suggests a high amount of overlap in the topic
 132 content between support and oppose responses.

143 We can directly compare the response-type conditioned posteriors of the combined model at any
 154 given K . In contrast, comparison of posteriors for the response-type specific models should account
 165 for the fact that the likelihood will in general suggest distinct K for each. With a prior in hand, we
 176 could simply marginalize over K in the posterior for each. We do not have such a posterior in hand
 187 unfortunately. A natural alternative is simply to compare models at the K at which each achieves the
 198 respective maximum in the likelihood. We highlight these maximum comparisons when plotting results
 209 over K .

210

221

222

223

224

225

226

227

228

229

230

231

232

233

234

235

236

237

238

239

240

241

242

243

244

245

246

247

248

249

250

251

252

253

254

255

256

257

258

259

260

261

262

263

264

265

266

267

268

269

270

271

272

273

274

275

276

277

278

279

280

281

282

283

284

285

286

287

288

289

290

291

292

293

294

295

296

297

298

299

300

301

302

303

304

305

306

307

308

309

310

311

312

313

314

315

316

317

318

319

320

321

322

323

324

325

326

327

328

329

330

331

332

333

334

335

336

337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384

385

386

387

388

389

390

391

392

393

394

395

396

397

398

399

400

401

402

403

404

405

406

407

408

409

410

411

412

413

414

415

416

417

418

419

420

421

422

423

424

425

426

427

428

429

430

431

432

433

434

435

436

437

438

439

440

441

442

443

444

445

446

447

448

449

450

451

452

453

454

455

456

457

458

459

460

461

462

463

464

465

466

467

468

469

470

471

472

473

474

475

476

477

478

479

480

481

482

483

484

485

486

487

488

489

490

491

492

493

494

495

496

497

498

499

500

501

502

503

504

505

506

507

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

569

570

571

572

573

574

575

576

577

578

579

580

581

582

583

584

8 *Puelma Touzel et al*

Table 1. Four low-order statistics of a data cloud in a $(K-1)$ -dimensional space ($\mu \in \mathbb{R}^{K-1}$) with mean position $\bar{\mu}$ and covariance matrix Σ_μ , having eigenvalues $\{\lambda_i\}_{i=1}^{K-1}$.

Data Cloud Property	Interpretation	Normalized Estimator
Position 	Specificity	$\frac{1}{\sqrt{K-1}} \mathbb{E} [\boldsymbol{\mu}]$, average distance from the origin (here: equal usage)
Size 	Variability	$\frac{1}{K-1} \sqrt{\sum_{i=1}^{K-1} \lambda_i}$, $\sqrt{\text{total variance}}$
Eccentricity 	(Lack of) Expressivity	$\frac{1}{K-1} (\sum_{i=1}^{K-1} \lambda_i)^2 / (\sum_{i=1}^{K-1} \lambda_i^2)$, intrinsic dimension
(+ve/-ve) Orientedness 	Alignment	$\frac{2}{(K-1)(K-2)} \sum_{\rho_{ij} \geq 0, i < j} \rho_{ij}$, sum of signed correlation, $(\rho_{ij} = \frac{(\Sigma_{\mu})_{ij}}{\sigma_i \sigma_j})$

2.7. Statistical measures of topic mixture statistics

Our primary interest is in the statistical structure of the posteriors over topic mixture weight vectors, θ eq. (1). However, the geometry of a domain can have a strong effect on the estimation of statistical properties and estimation using the θ -representation is biased because of the normalization constraint on weight vectors. A now well-established approach to performing statistical analysis on variables defined on the simplex (Pawlowsky-Glahn and Egozcue, 2001; Aitchison, 1982) maps these variables to the space of logarithms of relative weights (relative to the geometric mean, $g(\theta) = \prod_{i=1}^K \theta_i^{1/K}$). This bijective transformation maps the simplex to \mathbb{R}^{K-1} such that LogisticNormal distributions map back into their Normal counterparts. Applying this transformation for the computed models, we find that the average distance between all pairs from the set of estimated means $\hat{\mu}_d = \hat{\Gamma}\hat{x}_d$ over the three datasets is typically much larger than the size of the variation arising from the fitted covariance matrix parameter $\hat{\Sigma}$ (more than 95% of pairwise distances were larger than the total variance/dimension for 3 of the 4 comparisons (eqs. (3) and (4)) with the remainder at 85%; See Supplemental Material for details). Thus, the global geometry of the distribution is well-captured by the distribution of these means. We then focus on the empirical distribution of over $\mu = (\mu_1, \dots, \mu_{K-1})$, i.e. the set $\{\hat{\mu}_d\}_{d=1}^D$, one for each response in the response type dataset and clustering according to demographic statistics.

We developed and used four intuitive characteristics of a data cloud (see table 1). The first is simply the mean position $\bar{\mu}$, while the remaining three rely on the covariance, Σ_{μ} , estimated using the standard estimator $\hat{\Sigma}_{\mu} = \frac{1}{D-1} \sum_{d=1}^D (\hat{\mu}_d - \hat{\mu})(\hat{\mu}_d - \hat{\mu})^\top$, where $\hat{\mu} = \frac{1}{D} \sum_{d=1}^D \hat{\mu}_d$ is the mean estimate. In order to plot results across K without pure dimension effects displaying prominently, we apply appropriate normalization to each measure that gives them constant scaling in K . We label each measure according to its interpretation as a measure of ideology as follows. The position relative to equal usage (the origin in μ -space) measures how specific the usage is across topics. The size measures how much variability there is in how different individuals discuss that ideology. Eccentricity measures how the correlations limit the expressivity in how the ideology mixes topics. We measure this via reductions in a natural global measure of intrinsic dimension (Recanatesi et al., 2022). Finally, orientedness measures how strongly aligned or anti-aligned pairs of topics are.

Each measure is largely independent from the others and together the set is a largely comprehensive representation of unimodal, non-skewed data, since it covers global statistical features up to second-order. We discuss limitations and missing features of this set of measures in the discussion (section 4).

3. Results

3.1. Ideology and carbon tax support/opposition

To motivate our study of open-ended responses, we first show how support for the carbon tax depends on the demographic characteristics of the study's respondents. We find that support for carbon tax transitions from a majority to a minority opinion consistently across ideology-related feature dimensions

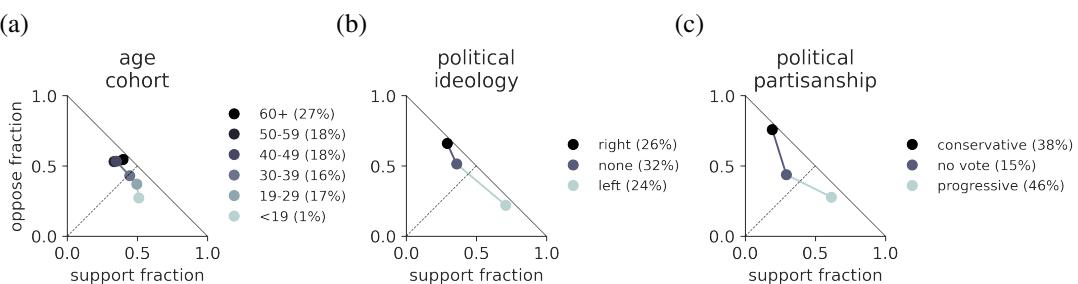


Figure 2. Carbon tax opposition vs. support varies strongly with political demographics. Fraction of respondents that stated they opposed vs. the fraction who stated they support the carbon tax over (a) age cohort, (b) political ideology, and (c) political partisanship (measured as voting preference; see Methods for details;). Addition of the remaining "not sure" fraction (not shown) sums to 1. Opposition is a majority (minority) opinion for dots above (below) the dashed line.

of left-to-right political ideology and liberal-to-conservative partisanship (see fig. 2). These transitions are relatively strong (they extend further away from equal support and opposition) in comparison with a transition with age, shown in fig. 2a as a baseline.

3.2. Discriminability of the word-level representation

Moving onto analysis of the open-ended responses, we first present the response type-conditioned word statistics in fig. 3a. One salient feature is how much the word ‘tax’ dominates in frequency within oppose responses, as compared with the support responses that recruit many frequent words. We then asked to what degree word content (which words appear and how many times) could be used to distinguish support and oppose responses. We performed logistic regression on the responses to distinguish support from oppose response type, finding best test accuracy around 80% (fig. 3a; see section 2 for details about the classification procedure). Word-level features are thus moderately informative in distinguishing support from opposition. This word-level information is spread heterogeneously over the vocabulary. In particular, by rerunning the classification on the n most contributing words (with contribution by frequency-weighted effect size) to each of the two target labels (support/oppose), we show in fig. 3b that test accuracy is maximized for about 100 words for each label (200 words overall) out of the more than 2500 stem words appearing in the data. We also show word clouds of the frequency-weighted effect size for $n = 100$ that display the most discriminating words. In sum, many words seem to contribute to a moderate level of discriminability.

A principal drawback of this word-level approach to assessing the discriminating power of the text responses is that individual words are assigned to one or the other response type. Yet, we know the words used by oppose and support responses overlap. For example, ‘carbon’ and ‘tax’ are used in high frequency by both response types. Removing this pair of words allows for higher discriminability (84% accuracy; not shown). That removing these central words increases discriminability suggests first-order word statistics are a poor representation of the semantics of the responses and that we should pursue approaches that employ latent representations that make use of words in different ways. In particular, it is the context in which these terms are used—what distinct ideas are evoked by the term—that ultimately distinguish the two response types and serves as a better semantic representation.

3.3. Topic modelling

To investigate the contextual origin of response type discriminability using the open-ended responses, we used a generative topic model with latent topic structure. We analyzed two model settings: (1)

10 Puelma Touzel et al.

1
201
302
403
504
605
706
807
908
1009
110
121
132
143
154
165
176
187
198
209
210
221
232
243
254
265
276
287
298
309
310
321
332
343
354
365
376
387
398
409
410
421
432
443
454
465
476
487
498
509
510
521
53
54
55

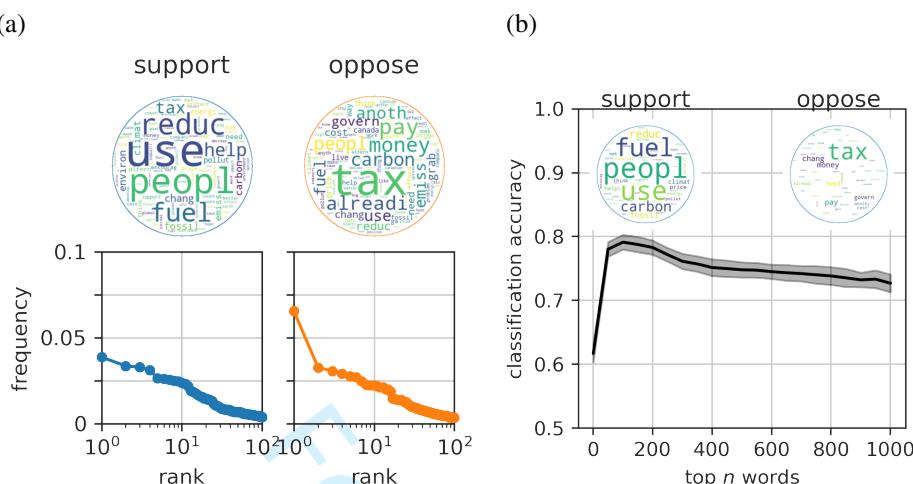


Figure 3. Word-level representations. (a) Tfifdf word cloud (top) and rank plot (bottom) for the word statistics of support and oppose response types. (b) Classification test accuracy of word-level content to predict response type (support/oppose; gray shows \pm standard deviation over 100 random train-test splits). Accuracy is plotted for projections of the data onto the n most predictive words (See Methods for details). Inset are word clouds of log-transformed frequency-weighted effect size for the $n = 100$ most predictive words for support (left) and oppose (right) classifications.

response type-specific models fit to responses of certain type (support/oppose) and (2) a single model fit to both types of responses combined.

To illustrate the topic content, in table 2, we show the top words associated with a given topic in a given model for the support and oppose response-specific models, as well as for the combined model (all for $K = 7$). A topic's top words can evoke an unambiguous semantic label for some topics, but not others. For example, topic a and e in the combined model evoke opposition because of the focus on cost of living and because of the word 'grab' as in 'tax grab', respectively. However, subjective inspection leaves many cases largely ambiguous as to whether a topic in the combined model is weighed more by support or oppose responses. Complementary to top words, top responses can also convey topic quality and even topic labels. For example, we perceived high semantic coherence in topic d, about unfairness of the tax on an international level (as well as, though to a lesser degree, in topic f about economic benefits, in topic c about its effectiveness in bringing about behaviour change, etc.). The latter were mostly not transparent in the list of top words, but only in the responses, i.e. in how these words were used. We list some top responses for these topics in the Supplemental Material for reference. It is hard to determine if the ambiguity in determining a topic label from a set of top responses arises from low topic quality reflecting a weakly clustered topic structure in the data or if it reflects limitations in our ability to identify salient topics. Another convoluting factor is that these topics are specific to the chosen value of K insofar as it imposes a specific number of topics, independent of any natural structure in data. For example, clusters of co-occurring words that are well-separated from each other will nevertheless likely be joined when K is too low and single clusters will likely be labeled incorrectly as multiple distinct clusters when K is too high.

To gain a deeper understanding of the relative changes in the topics as K increased without having to resort to *ad hoc* interpretation of the topics' top words, we developed two complementary analyses. First, we analyzed how topics breakup as K is increased (we call this a depth phylogeny). Inspired by methods of representing genetic phylogenies, we drew a tree structure obtained from connecting each topic to the topic in the model learned with 1 fewer topics that was nearest in the space of vocabulary word frequencies. In order to see how support and oppose responses spread their use of these topics

1
201 **Table 2.** Example topic content (for $K = 7$ and independent topic mixture prior, $\sigma = 1$).

Topic label	Frequency-ranked words
Support model	
a	chang, environ, climat, pollut, pay, mak, effect, protect, product, earth
b	peopl, help, energi, think, price, better, like, car, decrea, transport
c	emiss, less, increa, caus, govern, anoth, save, solut, drive, option
d	fuel, reduc, fossil, compani, consumpt, must, warm, global, incent, renew
e	tax, money, support, put, good, planet, believ, may, make, work
f	use, need, way, encourag, cost, hope, someth, gener, futur, thing
g	carbon, altern, sourc, find, get, develop, look, forc, creat, corpor
Oppose model	
a	use, fuel, work, car, altern, price, electr, like, stop, better
b	emiss, canada, countri, get, make, world, pollut, less, thing, big
c	carbon, alreadi, live, price, heat, caus, global, high, economi, warm
d	chang, reduc, anoth, grab, climat, noth, effect, believ, hurt, spend
e	tax, much, consum, problem, compani, industri, pocket, year, back, person
f	peopl, money, pay, fuel, need, cost, fossil, way, enough, mak
g	govern, put, think, environ, canadian, drive, emiss, incom, mani, averag
Combined model	
a	pay, alreadi, cost, enough, peopl, know, live, much, work, afford
b	money, mak, way, think, get, use, make, consum, environ, altern
c	fuel, reduc, carbon, chang, use, emiss, fossil, peopl, climat, help
d	canada, carbon, countri, effect, world, problem, pollut, believ, one, heat
e	tax, govern, anoth, grab, noth, put, price, car, canadian, still
f	need, like, support, thing, energi, someth, option, wast, resourc, gener
g	environ, tax, peopl, less, energi, pollut, encourag, altern, use, compani

287 for topics learned on both responses, we computed a topic depth phylogency for the combined model
 288 (see fig. 4a). Branch thickness denotes closeness³. We also added the relative support and oppose con-
 289 tribution to the mixture weight component of that topic as a piechart at each topic node in the tree. This
 290 allowed us to observe the cleaving of more coarse topics at low values of K into their finer-grained internal
 291 components as K increased. Reminiscent of phylogenetic analyses, we define a semantic-agnostic
 292 topic nomenclature using the row and column labels, respectively. In most cases, we see that topics are
 293 recruited by both response types so that single topics alone are insufficient for discriminating oppose
 294 from support responses. It is the weighting across multiple topics that gives sufficient information.
 295 That said, one salient feature of this particular tree is that topic 7f is the common ancestor for many of
 296 the topics at larger values of K . It predominantly loads onto support responses, recapitulating existing
 297 results in the literature of a diverse and diffuse set of topics discussed by those who support a carbon
 298 tax. In the following section, we will see that this $K = 7$ model is highly discriminative and this topic
 299 in particular conferred more than 99% accuracy.

400 In a second analysis, we assessed the quality of the topics over different K using the standard pair
 401 of metrics of exclusivity and semantic coherence (see Methods section). Running the model on each
 402 response type separately produced topics whose quality (fig. 4b) show that topic-averaged values give
 403 a linear trade-off between the two, with topic number K setting where along the trade-off the model
 404 resides. This suggests the linear combination as a family of scalar metrics of topic quality. Across
 405 different response types and values of K , the fixed sum (*i.e.* equal weights) of exclusivity and semantic
 406 coherence is highest in *oppose* responses, followed by *support* and *not sure*. We attribute this ordering
 407 to the singular focus that *oppose* responses have on the word ‘tax’, as compared to the much more
 408 diffuse responses (in word use) of *support* responses. Note, however, that the topic variance in each of
 409 the two quality metrics is at least as large as the difference in topic means at fixed topic number so the
 410

521 ³We use the inverse Euclidean distance between topics normalized by the sum of these distances over topics in a given model.

12 Puelma Touzel et al.

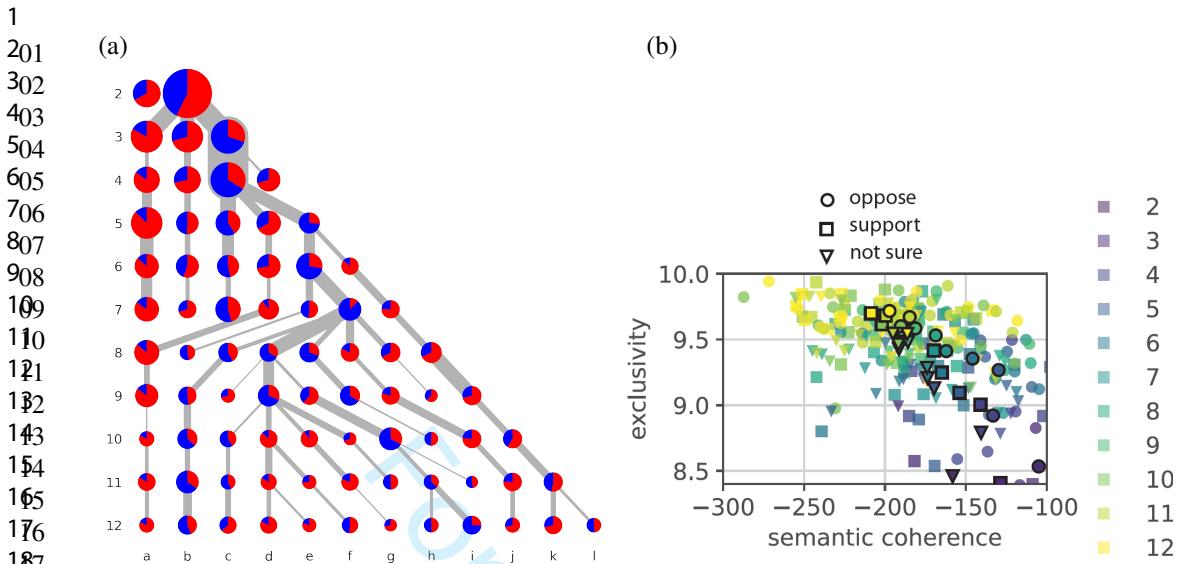


Figure 4. Topic structure over topic number, K . (a) Topic trees from varying the number of topics, K , in the combined model. A row corresponds to a given value of K starting at the top with $K = 2$ down to $K = 12$. The relative support (blue) and oppose (red) contribution to each response-averaged mixture weight component is shown as a piechart. Link thickness is inverse Euclidean distance between topics. (b) Topic quality for response type models. Topic-averaged exclusivity plotted against topic-averaged semantic coherence for topic number $K = 2, \dots, 12$ (colors) for the three response types (marker type). Black outlined markers denote the unweighted average over K .

aforementioned ranking is only apparent after averaging over topics. Semantic coherence bottomed out at around $K = 10$ topics. Topic quality for the composite model was qualitatively similar.

How predictive of response type are the way topics are recruited in the responses? In fig. 5a, we plot the empirical distribution of mixture weights, projected into its first two principal components, and colored coded with each response type. High separability (especially for $K = 7$) is observed across different topic number K , suggesting high discriminability. We performed a classification analysis on the support/oppose label similar to that using word use and presented at the beginning of our results (see fig. 3(b)), but now instead using the inferred topic mixture weights. Note that this representation is able to recruit information in the correlations among words that distinguishes the two response types. In fig. 5(b), we again show performance as a function of using the most predictive features (here not words, but topic mixture weights). Consistent with the high observed separability, we find accuracies approaching and more than 99% across all values of $K > 2$, demonstrating the power of this latent structure to predict support or opposition to the tax. Accuracies were only marginally lower at 95% or more for $\sigma = 0$ (not shown). For $K = 7$, near 100% accuracy is possible even using only a single topic and perfect accuracy for only two topics. As an example of more fine-grained analysis made possible by our approach, we focus in on these two topics. In fig. 5(c), we show the projection of the set of best-estimate mixtures across responses, colored by response type, into the plane in mixture space spanned by the components associated with these two topics. We see that the perfect classification results from topic 7 (topic 5) being strongly suppressed in support (oppose) responses. We further assessed this association by inspecting top responses for each topic (one from each is shown in fig. 5(c)). The responses in topic 7 describe the tax as a malicious injustice imposed on citizens, while those of topic 5 describe the effectiveness of the monetary incentive. We note that while the term ‘tax grab’ is the core idea in Topic 7 (‘grab’ is in 8.5% of oppose responses compared to only 1.5% of support responses), the idea recruits a much wider set of terms operationalized as the support of topic 7 on the

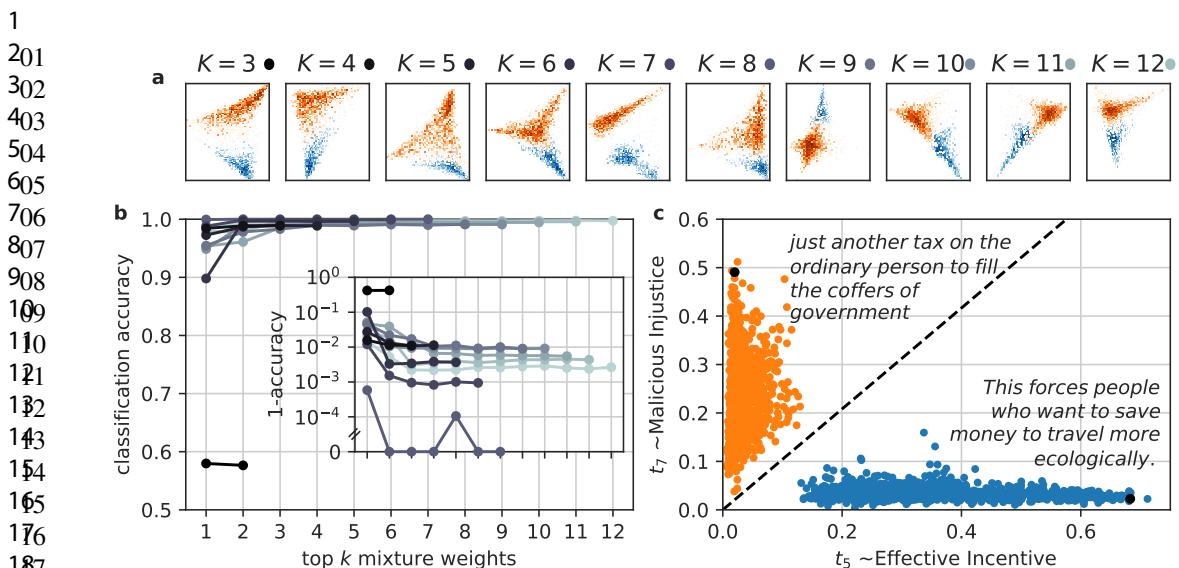


Figure 5. Topic-level Representations. (a) Projection of inferred mixture weights from the combined model onto the first two PCA components for a range of values of K . A histogram of counts on a fine-grained grid is shown using a white-to-blue and white-to-red color scale for support and oppose responses, respectively. (b) Classification test accuracy for using the best k of the mixture weights ranked by largest frequency-averaged effect size (black to gray is different values of K from 2 to 12). (c) The mixture estimates projected into the two top topics associated with the top $k = 2$ mixture weights for $K = 7$. The blacked-dashed line is the optimal classifier that achieves 100% accuracy. The text and mixture position of a top response for each topic is shown as a block dot.

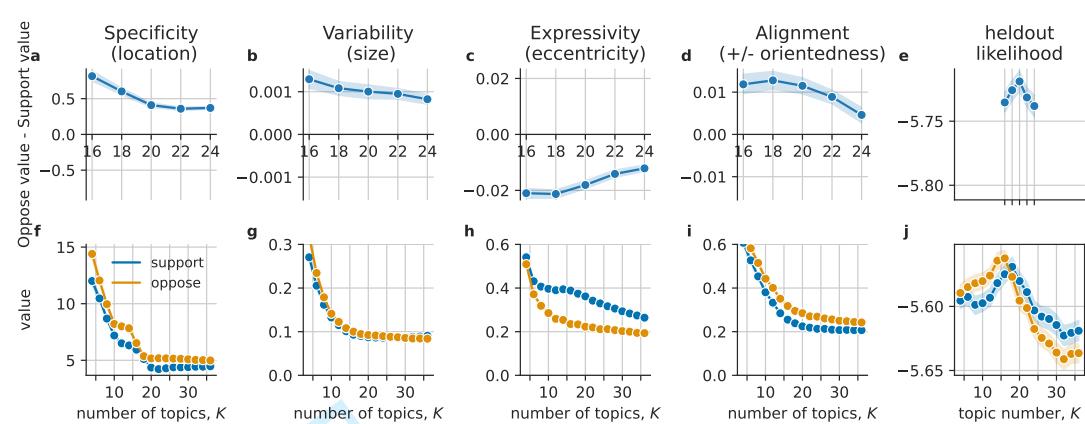
vocabulary. It is this wider set of words that confers the discriminability (in this case terms that are not used in high frequency by support responses).

We now turn to the topic-topic correlations to better understand the origin of the topic representation's higher classification performance over word frequency representations. We also arrive at analyses appropriate to investigating the structure of the mixture statistics that putatively hold distinctive features of the underlying ideology. We used the four measures of data cloud geometry summarized in table 1 to assess the topic-topic correlation structure using μ -representation as explained in the Methods. For the combined response type models, we can directly compare the support and oppose responses at each value of K , so we represent the results as the difference in the values of each measure over the two response types for a range of K around the maximum in the heldout Likelihood (positive values indicate oppose responses exceed support responses on that metric; fig. 6a-d). For the separate response type models, for which the best fitting value of K differs, we simply report both measures over a larger range of K (fig. 6f-i). By showing the heldout likelihood (fig. 6e,j), a comparison at respective maxima could be made. We find consistent results for both comparisons. Namely, oppose responses are more specific, somewhat more variable, much less expressive, and more aligned across a wide range of K around where the held-out likelihood is maximized. All three of the non-weak results are consistent with our hypothesis of a rigid ideology underlying carbon tax opposition.

4. Discussion

In this study, we presented a set of principled, quantitative measures and analyses to pin down topic structure in the learned parameters of the STM. We applied them to understand the effects of ideology behind carbon tax opinion in Canada. We find topic mixture weights derived from the learned models

14 Puelma Touzel et al.



165 **Figure 6. Evaluations of Statistical Measures of Mixture Statistics.** Top (a-d): The average difference
166 of the oppose value and the support value in each respective measure from the combined model (eq. (4)).
167 Bottom (f-j): The value of the measure for each of the support (blue) and oppose (orange) models
168 (eq. (3)). Lines are mean estimates and error bars are standard deviation of 100 (200) posterior samples
169 for top (bottom). Respective held-out model likelihoods are shown in (e) and (j).

222 are highly predictive of opposition to or support of the carbon tax and we presented a topic tree analysis
223 similar to phylogenetics to show that this performance arises when the model has a sufficient number
224 of topics that together are discriminating on each response. Finally, we proposed a set of statistical
225 measures on topic mixtures and evaluated them to find that the oppose responses had higher levels of
226 specificity and orientedness, and were less expressive. This suggests carbon tax opposition is explained
227 by its proponents through a more well-worn, coherent ideology. Together, these results on carbon tax
228 are thus suggestive of our initial hypothesis.

229 We anticipate the methods we present here can be useful to social scientists who are interested in
230 inferring social and political phenomenon from open-ended survey responses. Most obviously,
231 our methods could be used to further quantify the structure of collective beliefs. In a longitudinal
232 approach, they could also reveal the dynamics of these collective beliefs, and in particular their response
233 characteristics, for example from large-scale interventions or events (e.g. the COVID19 pandemic,
234 Russia-Ukraine war, etc.). A study of the timescales over which the effects of interventions decay
235 would inform the recent work on the notion of inoculation to misinformation. To push the metaphor,
236 the characteristics of the memory it induces have not yet been deeply probed (Roozenbeek et al., 2022).
237 Our work suggests that the topic neighborhood around the focus of any given intervention may play a
238 role via the topic-topic correlation structure. This could also impact policy design, for example in how
239 the carbon tax is promoted. Typical communications advice is to focus on a single issue/message to
240 not overload the target audience. However, neighborhood effects that are strongly restoring (Hall and
241 Bialek, 2019) could rapidly erase the effects of single issue interventions. Instead, a communications
242 strategy aimed at a local neighborhood of target issues in some appropriately defined issues space may
243 be more effective at loosening the grip that ideology has on opinion.

244 There are many potential follow-up areas to take this work. For one, the existing set of metrics of
245 mixture statistics only cover first and second order statistical moments. Moreover, there even aggregate
246 low order statistics missing from this set, e.g. how concentrated is the distribution around its
247 center (the expectation of $\|\mu\|$ might provide signal into how hollow is the data cloud). In another
248 direction, the efficiency of the topic representations suggests it could serve as an input embedding
249 space for use in deep learning approaches to large-scale social computations (e.g. those run on large
250 social media platforms). For example, recent work on reinforcement learning with human feedback for
251 large language models has focussed on generating consensus statements from a set of human-generated
252 messages (Huang et al., 2022).

1

201 opinions (Bakker et al., 2022). A currently limiting constraint of this approach is how many human-
302 generated text inputs can be used. A topic space learned from many opinions could circumscribe this
403 constraint.

504 We probed model behaviour over a range of K around the maximum in held-out model likelihood.
605 We were able to make conclusions from the fact that the sign of the difference of the values of the
706 metrics was unchanged over these ranges. This will not be true in general, in which case one approach
807 is to marginalize K out of the problem. In a fully Bayesian approach, given a prior over K , one would
908 infer a posterior over K and take the posterior average over any K -dependent quantities, e.g. average
109 topic quality or the statistical metrics. Given that models with different K have a different number of
110 parameters, a complexity-penalty term such as the AIC could be added to the posterior to account for
121 variable model complexity.

132 One direction we did not pursue here was exploring how different meta-data models affect the
143 results. Are there specific factor models of fewer covariates that best explain the data? As we moti-
154 vated our choice of the all-in model, this is a delicate analysis requiring disentangling various cases.
165 For example, is car use a confounder or collider of the effect of residence environment? Such distinc-
176 tions will influence estimated effect sizes (Wysocki et al., 2022). A first step along this direction is to
187 further analyze the results in the case of a sparsity prior on Γ . Including meta-data model dependence
198 in topic content is a related direction for future work. Note that in absence of rich metadata, the STM
209 model reduces roughly to the Correlated Topic Model (Blei and Lafferty, 2005), in which case topic
220 correlations can be studied via the mixture covariance matrix parameter, Σ_θ .

221 Finally, there is a broader question of the validity of the generative model class. In particular, are top-
232 ics models with latent structure suitable mathematical representations of semantic content of networks
243 of beliefs? How much is lost when representing language as text, without syntax and without positional
254 encoding? This question broaches broader philosophical issues around meaning in natural human lan-
265 guage. Given the recent evidence of (position-encoding) large language models that achieve in-context
276 learning (c.f. Mollo and Millière (2023)), we believe text is a sufficiently rich representation of natu-
287 ral language from which to extract semantics. Stripping text of its positional encoding must result in
298 some degree of loss of meaning. Topic models are nevertheless highly interpretable and we believe can
309 serve as useful complements to less easily interpretable large language models. While we often think
320 of syntax and word order as crucial to meaning, its importance to ecological communication may be
332 overstated. That said, ecological communication also presents challenges to using text to infer intended
343 meaning: a word can have one of many different meanings depending on who is speaking (Marti et al.,
2023).

354 In sum, we see a bright future for topic models in quantitative social science and beyond. We hope
365 that this work inspires their development, refinement and application as a research direction along
376 which more computational data science approaches contribute to the social aspects of climate change.
387

398 **Acknowledgments.** MPT would like to acknowledge helpful discussions with Matto Mildenberger, Kathryn Harrison, and
409 Dhanya Sridhar.

440 **Funding Statement.** MPT was funded under the CERC for Autonomous AI awarded to Irina Rish.

432 **Competing Interests.** The authors declare none.

443 **Data Availability Statement.** All supporting data was obtained from an existing Harvard Dataverse replication archive at
454 <https://doi.org/10.7910/DVN/3WBCH9>.

465 **Ethical Standards.** The research meets all ethical guidelines, including adherence to the legal requirements of the study country.

476 **Author Contributions.** MPT: Conceptualization (Equal), Formal analysis (Lead), Investigation (Lead), Methodology (Lead),
487 Software (Lead), Writing - original draft (Lead), Writing - review & editing (Lead). EL: Conceptualization (Equal), Data curation
498 (Lead), Supervision (Lead), Writing – review & editing (Supporting).

509 **Supplementary Material.** Supplementary Materials available online at <https://doi.org/10.7910/DVN/3WBCH9>

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

569

570

571

572

573

574

575

576

577

578

579

580

581

582

583

584

585

586

587

588

589

590

591

592

593

594

595

596

597

598

599

600

601

602

603

604

605

606

607

608

609

610

611

612

613

614

615

616

617

618

619

620

621

622

623

624

625

626

627

628

629

630

631

632

633

634

635

636

637

638

639

640

641

642

643

644

645

646

647

648

649

650

651

652

653

654

655

656

657

658

659

660

661

662

663

664

665

666

667

668

669

670

671

672

673

674

675

676

677

678

679

680

681

682

683

684

685

686

687

688

689

690

691

692

693

694

695

696

697

698

699

700

701

702

703

704

705

706

707

708

709

710

711

712

713

714

715

716

717

718

719

720

721

722

723

724

725

726

727

728

729

730

731

732

733

734

735

736

737

738

739

740

741

742

743

744

745

746

747

748

749

750

751

752

753

754

755

756

757

758

759

760

761

762

763

764

765

766

767

768

769

770

771

772

773

774

775

776

777

778

779

780

781

782

783

784

785

786

787

788

789

790

791

792

793

794

795

796

797

798

799

800

801

802

803

804

805

806

807

808

809

810

811

812

<p

1
201
References

- 302 J. Aitchison. The statistical analysis of compositional data. *Journal of the Royal Statistical Society. Series B (Methodological)*,
403 44(2):139–177, 1982. ISSN 00359246. URL <http://www.jstor.org/stable/2345821>.
- 504 J. E. Aldy and R. N. Stavins. The promise and problems of pricing carbon: Theory and experience. *The Journal of Environment
605 & Development*, 21(2):152–180, 2012.
- 706 M. A. Bakker, M. J. Chadwick, H. R. Sheahan, M. H. Tessler, L. Campbell-Gillingham, J. Balaguer, N. McAleese, A. Glaese,
807 J. Aslanides, M. M. Botvinick, and C. Summerfield. Fine-tuning language models to find agreement among humans with
diverse preferences, 2022. URL <https://arxiv.org/abs/2211.15006>.
- 908 D. Baldassarri and A. Gelman. Partisans without constraint: Political polarization and trends in american public opinion.
109 *American Journal of Sociology*, 114(2):408–446, 2008.
- 110 A. Baranzini, J. Goldemberg, and S. Speck. A future for carbon taxes. *Ecological Economics*, 32(3):395–412, 2000. ISSN
121 0921-8009. doi: [https://doi.org/10.1016/S0921-8009\(99\)00122-6](https://doi.org/10.1016/S0921-8009(99)00122-6). URL <https://www.sciencedirect.com/science/article/pii/S0921800999001226>.
- 132 S. Birch. Political polarization and environmental attitudes: A cross-national analysis. *Environmental Politics*, 29(4):697–718,
143 2020.
- 154 D. M. Blei and J. D. Lafferty. Correlated topic models. In *Proceedings of the 18th International Conference on Neural
165 Information Processing Systems*, NIPS'05, page 147–154, Cambridge, MA, USA, 2005. MIT Press.
- 176 R. J. Dalton and M. P. Wattenberg. *Parties Without Partisans: Political Change in Advanced Industrial Democracies*. Oxford
187 University Press, 03 2002. ISBN 9780199253098. doi: 10.1093/0199253099.001.0001. URL <https://doi.org/10.1093/0199253099.001.0001>.
- 198 S. Drews and J. C. Van den Bergh. What explains public support for climate policies? a review of empirical and experimental
209 studies. *Climate policy*, 16(7):855–876, 2016.
- 220 J. Eisenstein, A. Ahmed, and E. P. Xing. Sparse additive generative models of text. In *Proceedings of the 28th international
231 conference on machine learning (ICML-11)*, pages 1041–1048, 2011.
- 243 S. J. Gershman. How to never be wrong. *Psychonomic Bulletin and Review*, 26(1):13–28, 2019. ISSN 15315320. doi: 10.3758/
252 s13423-018-1488-8.
- 265 J. Grimmer, M. Roberts, and B. Stewart. *Text as Data: A New Framework for Machine Learning and the Social Sciences*.
274 Princeton University Press, 2022. ISBN 9780691207544. URL <https://books.google.ca/books?id=cL40EAAAQBAJ>.
- 287 E. Haites. Carbon taxes and greenhouse gas emissions trading systems: what have we learned? *Climate policy*, 18(8):955–966,
296 2018.
- 309 G. Hall and W. Bialek. The statistical mechanics of twitter communities. *Journal of Statistical Mechanics: Theory and
Experiment*, 2019(9):093406, sep 2019. doi: 10.1088/1742-5468/ab3af0. URL <https://dx.doi.org/10.1088/1742-5468/ab3af0>.
- 321 N. P. Kalmoe. Uses and abuses of ideology in political psychology. *Political Psychology*, 41(4):771–793, 2020.
- 333 D. Klenert, L. Mattauch, E. Combet, O. Edenhofer, C. Hepburn, R. Rafaty, and N. Stern. Making carbon pricing work for
340 citizens. *Nature Climate Change*, 8(8):669–677, Aug 2018. ISSN 1758-6798. doi: 10.1038/s41558-018-0201-2. URL
<https://doi.org/10.1038/s41558-018-0201-2>.
- 354 E. Lachapelle. Communicating about carbon taxes and emissions trading programs. In *Oxford research encyclopedia of climate
363 science*. Oxford University Press, 2017.
- 376 G. Lakoff. Why it matters how we frame the environment. *Environmental Communication*, 4(1):70–81, 2010. doi: 10.1080/
385 17524030903529749. URL <https://doi.org/10.1080/17524030903529749>.
- 398 L. Martí, S. Wu, S. T. Piantadosi, and C. Kidd. Latent Diversity in Human Concepts. *Open Mind*, 7:79–92, 03 2023. ISSN
409 2470-2986. doi: 10.1162/opmi_a_00072. URL https://doi.org/10.1162/opmi_a_00072.
- 421 A. M. McCright and R. E. Dunlap. The politicization of climate change and polarization in the american public's views of global
432 warming, 2001–2010. *The Sociological Quarterly*, 52(2):155–194, 2011.
- 443 A. M. McCright, R. E. Dunlap, and S. T. Marquart-Pyatt. Political ideology and views about climate change in the european
454 union. *Environmental Politics*, 25(2):338–358, 2016.
- 465 C. McGlade and P. Edkins. The geographical distribution of fossil fuels unused when limiting global warming to 2 c. *Nature*, 517
476 (7533):187–190, 2015.
- 487 M. Mildenberger, E. Lachapelle, K. Harrison, and I. Stadelmann-Steffen. Limited impacts of carbon tax rebate programmes
498 on public support for carbon pricing. *Nature Climate Change*, 12(2):141–147, 2022. ISSN 17586798. doi: 10.1038/
509 s41558-021-01268-3.
- 521 D. C. Mollo and R. Millière. The vector grounding problem, 2023.
- 533 V. Pawlowsky-Glahn and J. J. Egozcue. Geometric approach to statistical analysis on the simplex. *Stochastic Environmental
544 Research and Risk Assessment*, 15(5):384–398, Oct 2001. ISSN 1436-3259. doi: 10.1007/s004770100077. URL <https://doi.org/10.1007/s004770100077>.
- 555 M. Povitkina, S. Carlsson Jagers, S. Matti, and J. Martinsson. Why are carbon taxes unfair? Disentangling public perceptions of
566 fairness. *Global Environmental Change*, 70(September), 2021. ISSN 09593780. doi: 10.1016/j.gloenvcha.2021.102356.
- 577 B. G. Rabe. *Can we price carbon?* MIT Press, 2018.
- 589 S. Recanatesi, S. Bradde, V. Balasubramanian, N. A. Steinmetz, and E. Shea-Brown. A scale-dependent measure of system
599 dimensionality. *Patterns*, 3(8), Aug 2022. ISSN 2666-3899. doi: 10.1016/j.patter.2022.100555. URL <https://doi.org/10.1016/Cambridge University Press>

- 1
201 j.patter.2022.100555.
302 M. E. Roberts, B. M. Stewart, D. Tingley, C. Lucas, J. Leder-Luis, S. K. Gadarian, B. Albertson, and D. G. Rand. Structural topic
403 models for open-ended survey responses. *American Journal of Political Science*, 58(4):1064–1082, 2014. ISSN 15405907.
504 doi: 10.1111/ajps.12103. URL <http://dvn.iq.harvard.edu/dvn/dv/ajps>.
605 M. E. Roberts, B. M. Stewart, and D. Tingley. Stm: An R package for structural topic models. *Journal of Statistical Software*,
706 91:1–40, 2019.
807 J. Roozenbeek, S. van der Linden, B. Goldberg, S. Rathje, and S. Lewandowsky. Psychological inoculation improves resilience
908 against misinformation on social media. *Science Advances*, 8(34):eab06254, 2022. doi: 10.1126/sciadv.abo6254. URL
100 <https://www.science.org/doi/abs/10.1126/sciadv.abo6254>.
110 I. Savin, S. Drews, S. Maestre-Andrés, and J. van den Bergh. Public views on carbon taxation and its fairness: a computational-
120 linguistics analysis. *Climatic Change*, 162(4):2107–2138, 2020. ISSN 15731480. doi: 10.1007/s10584-020-02842-y.
130 E. Tvinnereim, K. Fløttum, Ø. Gjerstad, M. P. Johannesson, and Å. D. Nordø. Citizens' preferences for tackling climate change.
140 Quantitative and qualitative analyses of their freely formulated solutions. *Global Environmental Change*, 46(May):34–41,
150 2017. ISSN 09593780. doi: 10.1016/j.gloenvcha.2017.06.005.
160 D. Welsby, J. Price, S. Pye, and P. Ekins. Unextractable fossil fuels in a 1.5° c world. *Nature*, 597(7875):230–234, 2021.
170 A. C. Wysocki, K. M. Lawson, and M. Rhemtulla. Statistical control requires causal justification. *Advances in Methods and
180 Practices in Psychological Science*, 5(2):25152459221095823, 2022. doi: 10.1177/25152459221095823. URL <https://doi.org/10.1177/25152459221095823>.